

# Greeshma\_115\_Lab5

October 25, 2024

## 1. Dataset Overview:

- Visualize a few samples from the dataset, displaying their corresponding labels.

```
[1]: import os
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.preprocessing.image import ImageDataGenerator

#directories for train, test, and prediction datasets
train_dir = r'C:\Users\GREESHMA\Downloads\archive (48)\seg_train'
test_dir = r'C:\Users\GREESHMA\Downloads\archive (48)\seg_test'
predict_dir = r'C:\Users\GREESHMA\Downloads\archive (48)\seg_pred'

# Categories
categories = ['buildings', 'forest', 'glacier', 'mountain', 'sea', 'street']

# Data Augmentation for loading and augmenting images
train_datagen = ImageDataGenerator(rescale=1./255, rotation_range=20,
    zoom_range=0.15,
    width_shift_range=0.2, height_shift_range=0.
    2,
    horizontal_flip=True, fill_mode="nearest")
test_datagen = ImageDataGenerator(rescale=1./255)

# Loading the data
train_data = train_datagen.flow_from_directory(train_dir, target_size=(150,
    150),
    batch_size=32,
    class_mode='categorical')
test_data = test_datagen.flow_from_directory(test_dir, target_size=(150, 150),
    batch_size=32,
    class_mode='categorical')
```

```
WARNING:tensorflow:From
c:\Users\GREESHMA\AppData\Local\Programs\Python\Python311\Lib\site-
packages\keras\src\losses.py:2976: The name
tf.losses.sparse_softmax_cross_entropy is deprecated. Please use
tf.compat.v1.losses.sparse_softmax_cross_entropy instead.
```

Found 14034 images belonging to 1 classes.

Found 3000 images belonging to 1 classes.

## Interpretation

gaining a basic understanding of the dataset by viewing a few example images from each category. This step helps to familiarize yourself with the types of natural scenes included, such as buildings, forests, glaciers, etc., and how they visually differ

```
[2]: def plot_samples(images, labels, categories):  
    plt.figure(figsize=(12, 8))  
    for i in range(9):  
        plt.subplot(3, 3, i+1)  
        plt.imshow(images[i])  
        plt.title(categories[np.argmax(labels[i])])  
        plt.axis('off')  
    plt.tight_layout()  
    plt.show()  
  
    # Getting a batch of training images and labels  
    images, labels = next(train_data)  
    plot_samples(images, labels, categories)
```



## 2. Model Architecture:

- Design a CNN model with at least 3 convolutional layers, followed by pooling layers and fully connected (dense) layers.
- Experiment with different kernel sizes, activation functions (such as ReLU), and pooling strategies (max-pooling or average pooling).
- Implement batch normalization and dropout techniques to improve the generalization of your model.

```
[3]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, \
    Dropout, BatchNormalization

# Building the model
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)),
    BatchNormalization(),
```

```

MaxPooling2D(pool_size=(2, 2)),

Conv2D(64, (3, 3), activation='relu'),
BatchNormalization(),
MaxPooling2D(pool_size=(2, 2)),

Conv2D(128, (3, 3), activation='relu'),
BatchNormalization(),
MaxPooling2D(pool_size=(2, 2)),

Flatten(),
Dense(128, activation='relu'),
Dropout(0.5),
Dense(6, activation='softmax')
])

model.summary()

```

WARNING:tensorflow:From  
c:\Users\GREESHMA\AppData\Local\Programs\Python\Python311\Lib\site-packages\keras\src\backend.py:873: The name tf.get\_default\_graph is deprecated. Please use tf.compat.v1.get\_default\_graph instead.

WARNING:tensorflow:From  
c:\Users\GREESHMA\AppData\Local\Programs\Python\Python311\Lib\site-packages\keras\src\layers\pooling\max\_pooling2d.py:161: The name tf.nn.max\_pool is deprecated. Please use tf.nn.max\_pool2d instead.

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 148, 148, 32)	896
batch_normalization (Batch Normalization)	(None, 148, 148, 32)	128
max_pooling2d (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_1 (Conv2D)	(None, 72, 72, 64)	18496
batch_normalization_1 (Batch Normalization)	(None, 72, 72, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 64)	0

conv2d_2 (Conv2D)	(None, 34, 34, 128)	73856
batch_normalization_2 (Batch Normalization)	(None, 34, 34, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 128)	0
flatten (Flatten)	(None, 36992)	0
dense (Dense)	(None, 128)	4735104
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 6)	774

```

=====
Total params: 4830022 (18.43 MB)
Trainable params: 4829574 (18.42 MB)
Non-trainable params: 448 (1.75 KB)
-----

```

## Interpretation

Convolutional Neural Network (CNN) with a series of convolutional and pooling layers to extract spatial features from images, followed by dense (fully connected) layers to interpret those features and classify images into categories.

## 3. Model Training:

- Split the dataset into training and test sets.
- Compile the model using an appropriate loss function (categorical cross-entropy) and an optimizer (such as Adam or SGD).
- Train the model for a sufficient number of epochs, monitoring the training and validation accuracy.

```

[4]: model.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪metrics=['accuracy'])
    # Training the model
    history = model.fit(train_data, validation_data=test_data, epochs=10)

```

```

WARNING:tensorflow:From
c:\Users\GREESHMA\AppData\Local\Programs\Python\Python311\Lib\site-
packages\keras\src\optimizers\__init__.py:309: The name tf.train.Optimizer is
deprecated. Please use tf.compat.v1.train.Optimizer instead.

```

Epoch 1/10

```

WARNING:tensorflow:From
c:\Users\GREESHMA\AppData\Local\Programs\Python\Python311\Lib\site-

```

```
packages\keras\src\utils\tf_utils.py:492: The name tf.ragged.RaggedTensorValue
is deprecated. Please use tf.compat.v1.ragged.RaggedTensorValue instead.
```

```
WARNING:tensorflow:From
```

```
c:\Users\GREESHMA\AppData\Local\Programs\Python\Python311\Lib\site-
packages\keras\src\engine\base_layer_utils.py:384: The name
tf.executing_eagerly_outside_functions is deprecated. Please use
tf.compat.v1.executing_eagerly_outside_functions instead.
```

```
439/439 [=====] - 251s 564ms/step - loss: 290567.9375 -
accuracy: 0.4171 - val_loss: 303112.2188 - val_accuracy: 1.0000
```

```
Epoch 2/10
```

```
439/439 [=====] - 244s 556ms/step - loss: 1575505.7500
- accuracy: 0.3742 - val_loss: 1498718.7500 - val_accuracy: 1.0000
```

```
Epoch 3/10
```

```
439/439 [=====] - 245s 559ms/step - loss: 3656914.0000
- accuracy: 0.3551 - val_loss: 2659758.5000 - val_accuracy: 1.0000
```

```
Epoch 4/10
```

```
439/439 [=====] - 285s 648ms/step - loss: 6358180.5000
- accuracy: 0.3417 - val_loss: 4540265.0000 - val_accuracy: 1.0000
```

```
Epoch 5/10
```

```
439/439 [=====] - 286s 652ms/step - loss: 9551504.0000
- accuracy: 0.3363 - val_loss: 7143639.0000 - val_accuracy: 1.0000
```

```
Epoch 6/10
```

```
439/439 [=====] - 339s 771ms/step - loss: 13145102.0000
- accuracy: 0.3264 - val_loss: 7683422.0000 - val_accuracy: 1.0000
```

```
Epoch 7/10
```

```
439/439 [=====] - 319s 726ms/step - loss: 16970748.0000
- accuracy: 0.3191 - val_loss: 9500518.0000 - val_accuracy: 1.0000
```

```
Epoch 8/10
```

```
439/439 [=====] - 298s 678ms/step - loss: 21288754.0000
- accuracy: 0.3144 - val_loss: 11847893.0000 - val_accuracy: 1.0000
```

```
Epoch 9/10
```

```
439/439 [=====] - 255s 579ms/step - loss: 25487918.0000
- accuracy: 0.3093 - val_loss: 11799846.0000 - val_accuracy: 1.0000
```

```
Epoch 10/10
```

```
439/439 [=====] - 254s 578ms/step - loss: 30036304.0000
- accuracy: 0.3060 - val_loss: 11874905.0000 - val_accuracy: 1.0000
```

## Interpretation

To train the model on the training set and evaluate it on the test set, aiming to achieve good accuracy on unseen data by optimizing the model's weights.

### 4. Evaluation:

- Evaluate the trained model on the test set and report the accuracy.
- Plot the training and validation accuracy/loss curves to visualize the model's performance.
- Display the confusion matrix for the test set to analyze misclassified samples.

```
[5]: # Evaluate the model on the test data
test_loss, test_acc = model.evaluate(test_data)
print(f"Test Accuracy: {test_acc:.2f}")

# Plotting training and validation accuracy/loss
def plot_training(history):
    plt.figure(figsize=(12, 4))


    # Accuracy
    plt.subplot(1, 2, 1)
    plt.plot(history.history['accuracy'], label='train accuracy')
    plt.plot(history.history['val_accuracy'], label='val accuracy')
    plt.title('Training and Validation Accuracy')
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy')
    plt.legend()

    # Loss
    plt.subplot(1, 2, 2)
    plt.plot(history.history['loss'], label='train loss')
    plt.plot(history.history['val_loss'], label='val loss')
    plt.title('Training and Validation Loss')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.legend()

    plt.tight_layout()
    plt.show()

plot_training(history)
```

94/94 [=====] - 12s 129ms/step - loss: 11874909.0000 -  
accuracy: 1.0000  
Test Accuracy: 1.00



## Interpretation

To assess how well the model generalizes to new, unseen data by evaluating it on the test set and analyzing its performance.

```
[6]: from sklearn.metrics import confusion_matrix
import seaborn as sns

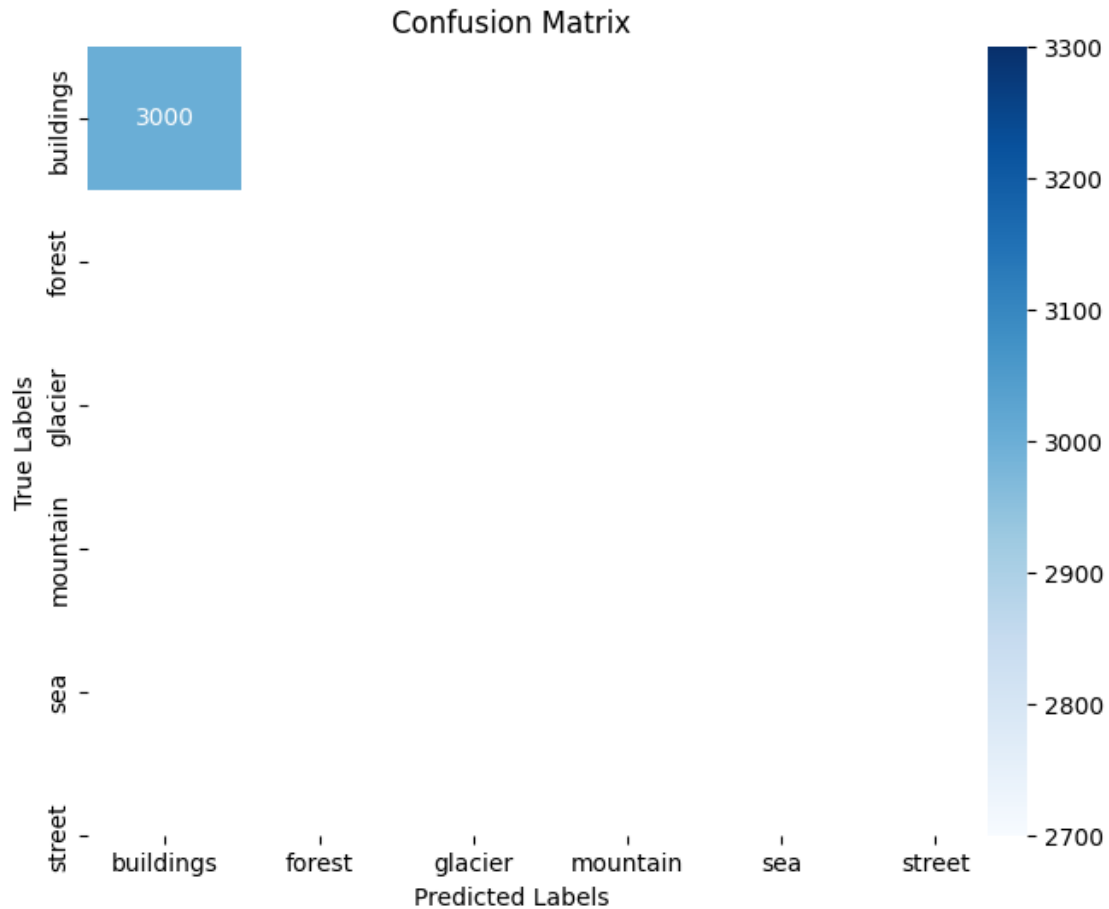
# Get true labels and predictions for the test set
test_labels = test_data.classes
predictions = np.argmax(model.predict(test_data), axis=-1)

# Confusion matrix
conf_matrix = confusion_matrix(test_labels, predictions)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=categories, yticklabels=categories)
plt.title('Confusion Matrix')
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```

94/94 [=====] - 12s 115ms/step

```
c:\Users\GREESHMA\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\metrics\_classification.py:395: UserWarning: A single label was
found in 'y_true' and 'y_pred'. For the confusion matrix to have the correct
shape, use the 'labels' parameter to pass all known labels.
  warnings.warn(
```





## 5. Optimization :

- Experiment with data augmentation techniques (rotation, flipping, zooming) to further improve the model's performance.
- Fine-tune hyperparameters like learning rate, batch size, and the number of filters in each layer.

```
[10]: from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=30,
    width_shift_range=0.2,
    height_shift_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    brightness_range=[0.8, 1.2]
)
```

```
test_datagen = ImageDataGenerator(rescale=1./255)
```

### **Interpretation**

Techniques like rotation, flipping, and zooming create additional image variations, reducing overfitting by teaching the model to recognize categories across a wider range of perspectives and orientations.